



Topological Navigation for Multi-Robot Pursuit-Evasion

Aren Petrossian^{1, 2}, Arsh Singhal^{1, 3}, & Joseph Hardin^{1, 4}

¹Georgia Institute of Technology



Introduction

Objective: Adapt topological navigation + effective RL agent for a multi-robot system to collaboratively locate and corner smart, mobile evader


Environment Constraints:

- Limited Sensing → No metric (SLAM) map
- Limited Bandwidth → Can't communicate large data
- Unknown environment → Have to explore
- Limited FOV and speed

Motivation: Efficient, fast localization of targets using a system of robots would be very useful in real-life deployments of indoor security robots → minimize risks for humans. Also, an explosion in real-life robotics, so it's important now. Our environment constraints were designed to closely mimic real-life difficulties.

Prior Work:

- Pursuit-evasion has been studied in many fields
- Other robotics work use known maps, single pursuer, centralized communication, etc.
- No use of topologies + RL (TopoNav) with multiple robots




Methods

Core Contributions & Insights:

- Realistic scenario needed for useful applications → **lack of intense computation or communication**
- Framework for **multi-robot topological mapping** + exploration through **shared landmarks**
- MDP logic + deep RL agent for **autonomous subgoal selection**
- Designed **extrinsic + intrinsic rewards** to balance different goals

POTENTIAL FIELD EVADER



Static: A* attraction + wall repulsive + wall raising
Dynamic: Repulsion radius near pursuers

Fallback: Pushed in opposite direction → tries using A* to get around pursuers and to goal

RL-AGENT PURSUER (Single Policy)

Markov Decision Process (MDP)

State: $G(V, E)$ of topological map
Edges: traversable path from nodes
Nodes: Pre-defined landmark
• # times visited • t since last visit
• who visited last • uncertainty
• dist. to curr loc • # of neighbors

Action: best subgoal g in V

Rewards: Extrinsic + Intrinsic

Extrinsic: Seeing Evader
Intrinsic:
• increase area explored
• re-visit stale (by time) nodes
• help expand frontier nodes

Loop till Evader found

Local Planner + Movement:
• Find path to g through graph
• Assume optimal local planning to get to neighbor node (A*)

While Moving (src → tgt):
• If landmark seen, add node to graph with edge to src and tgt

DQN: $Q(s_g, g; \theta)$
• GCN to get node embeddings
• Feed through Q network and pick subgoal with max Q

Reached node g :
• Update state and reward
• Expand neighbor nodes of g if seen by another pursuer
• If no unvisited neighbors, add 'artificial' node towards any unexplored region

Experimental Design

BASELINE METHODS

- Random Agent:** Picks random subgoal to traverse to as topological graph of landmarks grow normally
- Greedy Agent:** Greedily selects a seen, but unvisited, node as subgoal from topology
- Naive Patrol (extra knowledge):** Doesn't need to explore since it has hardcoded patrol paths for each pursuer. Once any pursuer finds evader, all pursuers towards evader with help of A* to avoid immediate obstacles

*If seen evader → chase evader, note close landmarks

METRICS

Run each pursuer strategy on varying evader speeds

- Full Capture %:** All 3 pursuers see evader
- # of Successful Pursuers:** 0-3
- % ≥ 1, ≥ 2, ≥ 3 Successful Pursuers**
- Average Episode Time:** Iterations to Chase/Evade




Figure 1: Hard-coded patrol paths for each pursuer

Training Reward:

$$R_{ext} = 50 \cdot \mathbf{1}_{\{\text{caught}\}}$$

$$R_{int} = 40 \times \frac{\Delta A}{4000} - 0.5 \frac{(\bar{n} - 1)}{\sqrt{\bar{n}}} + 0.1 \tanh\left(\frac{t}{200}\right)$$

ΔA : new area seen, \bar{n} : num visits, t : time since last visit

Hypotheses:

- H1:** RL-agent achieves higher capture rate than all baselines, esp. with faster evader
- H2:** Intrinsic reward will push RL-agent to explore a lot
- H3:** Naive patrol will be strong because of extra knowledge given to it

Results




Table 2: Full Capture Rate (%)

Strategy	0.33	0.50	0.75	0.90	1.0
Greedy	46	16	22	16	12
Naive Patrol	100	100	74	60	36
Random	30	22	10	10	2
RL-Agent	52	58	60	66	54

Table 1: Avg. Chase Time for Full Captures

Strategy	0.33	0.50	0.75	0.90	1.0
Greedy	577.3	510.0	377.0	848.6	350.5
Naive Patrol	100.4	94.5	109.6	94.0	73.2
Random	609.8	672.6	434.6	983.8	1488.0
RL-Agent	543.2	533.2	763.0	648.2	667.1

Table 3: Avg. Number of Successful Pursuers

Strategy	0.33	0.50	0.75	0.90	1.0
Greedy	2.28	1.52	1.45	1.20	1.04
Naive Patrol	3	3	2.38	1.80	1.78
Random	1.86	1.62	0.88	1.02	0.82
RL-Agent	2.52	2.36	2.34	2.32	2.30




Figure 2: Map with Pre-Set Landmarks




Figure 3: Evader Pot. Field from Fig. 1




Figure 4: Example Pursuer Topology

Takeaways:

- RL-Agent outperforms all baselines in key metrics besides against Naive Patrol for slower evaders (0.33 - 0.75 speed). Note that Naive Patrol is not a fair comparison since it has more hard-coded knowledge (i.e. patrol routes + uses A*)
- RL-Agent has 100% capture rate for evaders of all speeds.

Discussion of Results

Notable Phenomena:

- Robust Learned Subgoal Selection:** RL-Agent maintains ≥ 50% full-capture rate while other baselines drop to 10%-30%
- Strategic vs. Reactive Coverage:** Naive patrol collapses at top speed → simple compass-heading lacks flexibility




Figure 5: Learned Strategies → Cornering (Left) and Cutoff (Right)

- Efficiency of Topological Fusion:** RL approach matches/outperforms baseline that implicitly "sees" full grid. Lightweight topology framework is successful
- Faster Chase Time for Naive Patrol:** Metric maps has faster chase down since exact navigation details known
- Consistent Capture Rate (through speeds):** RL-agent builds performance as training speed increases → could use more episodes at each speed




Figure 6: Potential Field Evader + Naive Pursuers in Simulator (top) and Robotarium Deployment (bottom)

Limitations

- Limited Training + Experimentation:** Many more combinations of rewards, training schedules, features, etc. could be tested to reach optimality
- Not Robust to Randomness:** Evader has semi-random start but pursuers are fixed starts. Haven't fully tested or trained on random starts
- Not Robust to Evader Variety:** We fixed evader potential field and A* parameters for "realistic" evader, but could be smarter in real-life
- Pre-Set Landmarks:** Landmarks are static and pre-defined. No real-time detection or formation
- Suitability for Real Deployment:** Only worked with 2D environment, ignores struggles of 3D world. Also, loose "capturing" / apprehension mechanism

Conclusion

RL-Agent achieves **robust coordination, lightweight mapping**

- Reliably **outperforms other baselines** in capture rate, chase time

Builds on prior work in Pursuit-Evasion and navigation such as TopoNav

- First time **combining topologies, deep RL, multi-robots** for Pursuit-Evasion

Real-world application for 'apprehend-the-intruder' with robot security guards

Future Direction: Investigate more reward structures and features to enable faster chase-down. Generate landmarks real-time based on 3D features